


# Convex Optimization I

---

Using material from Stephen Boyd

# Why do we need optimization in robotics?


- Gives us a way to frame robotics problems mathematically
- VERY widely used
- Example: Inverse Optimal Control:



Learning Objective Functions for Manipulation  
[Kalakrishnan et al., ICRA 2013]

# Why do we need optimization in robotics?

- Example: Simultaneous Localization and Mapping (SLAM)



Keyframe-Based Visual-Inertial SLAM Using Nonlinear Optimization  
[Leutenegger et al., RSS 2013]

# Convex Optimization

- Convex optimization is a mature field with deep mathematical foundations
- It is so powerful that it's often worth it to
  - Work hard to reformulate your problem as convex
  - Approximate non-convex objective functions as convex
  - Use solution to approximation to start search for solution to the real problem
- It scales well with dimensionality
  - Convex optimization routinely solves problems with 1000s of variables
- Convex optimizers are fast (usually)

# Outline

- Calculus Review
- Convex Sets
- Convex Functions
- Unconstrained Optimization

# Set Notation

$$X = \{x \mid a^T x \leq b, x \in C, a \in \mathbf{R}^n\}$$

$X$  is the set of  $x$ s such that  $a^T x \leq b$  is true for  $x$  in the set  $C$  where  $a$  is a vector in a Euclidian space of dimension  $n$

# Review: Functions

- Functions are defined as:

$$f : A \rightarrow B$$


- “f maps elements in the set A to elements in the set B”
- The set  $A$  is the **domain** of f
- The set  $B$  is the **range** of f
- Example:

$$f : \mathbf{R}^n \rightarrow \mathbf{R}^m$$

“Function f maps n-dimensional vectors to some m-dimensional vectors”

# Review: Derivatives

- Derivatives can get complicated!
- Keep this in mind: A derivative is a linear approximation of how a function changes at a certain point



- The derivative of  $f(x)$  is the ratio between an infinitesimal change in an input variable  $x$  and the resulting change in the output  $f(x)$



## Review: Derivatives

- Recall the definition for a derivative  $f: \mathbf{R} \rightarrow \mathbf{R}$

$$Df(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- We can write a similar definition for  $f: \mathbf{R}^n \rightarrow \mathbf{R}^m$

# Review: Derivatives

- Suppose  $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$
- The function  $f$  is differentiable at  $x$  if there exists a matrix  $Df(x) \in \mathbb{R}^{m \times n}$  that satisfies

This is a matrix

$$\lim_{\substack{z \in \text{dom } f, z \neq x, z \rightarrow x}} \frac{\|f(z) - f(x) - Df(x)(z - x)\|_2}{\|z - x\|_2} = 0$$

- $Df(x)$  is called the derivative (or Jacobian) of the function
- $Df(x)$  can be computed by computing partial derivatives

$$Df(x)_{ij} = \frac{\partial f_i(x)}{\partial x_j}, \quad i = 1, \dots, m, \quad j = 1, \dots, n$$

# Review: Gradient

- When  $f$  is real-valued  $(i.e., f : \mathbf{R}^n \rightarrow \mathbf{R})$  the derivative  $Df(x)$  is a row vector (a  $1 \times n$  matrix)

Range must be 1-dimensional!

- The transpose of the derivative is the **gradient**:

$$\nabla f(x) = Df(x)^T$$

- Again, you can compute the gradient by taking partial derivatives:

$$\nabla f(x)_i = \frac{\partial f(x)}{\partial x_i}, \quad i = 1, \dots, n.$$

## Review: Second Derivative

- When  $f$  is real-valued  $(i.e., f : \mathbf{R}^n \rightarrow \mathbf{R})$  the **second derivative** is called the **Hessian Matrix**:  $\nabla^2 f(x)$

$$\nabla^2 f(x)_{ij} = \frac{\partial^2 f(x)}{\partial x_i \partial x_j}, \quad i = 1, \dots, n, \quad j = 1, \dots, n,$$

- Recall that the second derivative is the derivative of the first derivative:

$$D\nabla f(x) = \nabla^2 f(x)$$

# Questions


- Suppose we have a real-valued function  $f : \mathbb{R}^n \rightarrow \mathbb{R}$ 
  1. What are the dimensions of the gradient vector  $\nabla f(x)$ ?
  2. What are the dimensions of the Hessian matrix  $\nabla^2 f(x)$ ?

# Convex Sets


---

# Convex sets and functions

- Convexity is a restriction on shapes and functions
  - Convex optimization only works when everything is convex!
- We will cover definitions of convexity for shapes and functions
- You can use these to build convex sets/functions for the problems you care about



A convex set




A convex function

# Convex Sets

- Convex set: contains line segment between any two points in the set.  $C$  is a *convex set* if:

$$x_1, x_2 \in C, \quad 0 \leq \theta \leq 1 \quad \implies \quad \theta x_1 + (1 - \theta)x_2 \in C$$

- Examples:





# Important Types of Convex Sets: Hyperplane


- **Hyperplane**: A set of points that have a constant inner product with vector  $a$

$$\{x \mid a^T x = b\} \quad (a \neq 0)$$

same as  $a \cdot x$

- Another way to define it:

$$\{x \mid a^T (x - x_0) = 0\}$$




# Important Types of Convex Sets: Halfspace

- **Halfspace**: A hyperplane with an inequality

$$\{x \mid a^T x \leq b\} \quad (a \neq 0)$$


- Another way to define it :

$$\{x \mid a^T (x - x_0) \leq 0\}$$



# Important Types of Convex Sets: Polyhedron

- **Polyhedron**: The intersection of a finite number of halfspaces and hyperplanes



- Another way to define it: The set of solutions to a set of linear inequalities and equalities:

$$Ax \leq b$$

$$Cx = d$$

# Important Convexity-Preserving Operations on Sets

- **Intersection** preserves convexity


If  $S_1$  and  $S_2$  are convex, then  $S_1 \cap S_2$  is convex

- It follows that the intersection of any number of convex sets is convex
- **Affine functions** preserve convexity

$$f : \mathbf{R}^n \rightarrow \mathbf{R}^m \quad f(x) = Ax + b \text{ with } A \in \mathbf{R}^{m \times n}, b \in \mathbf{R}^m$$

- Examples of affine functions
  - Scaling
  - Translation
  - Projection

How do we know a polyhedron is always convex?



# Convex Functions

---

# Convex Functions

The domain of the function

$f : \mathbf{R}^n \rightarrow \mathbf{R}$  is convex if  $\mathbf{dom} f$  is a convex set and

$$f(\theta x + (1 - \theta)y) \leq \theta f(x) + (1 - \theta)f(y)$$




for all  $x, y \in \mathbf{dom} f$ ,  $0 \leq \theta \leq 1$



- I.e. the line segment between  $(x, f(x))$  and  $(y, f(y))$  lies above the graph of  $f$

# Advantage of convex functions


- Convex functions have only one local minimum!
  - That means local methods can find the global optimum!





# Concave Functions


- Concave functions are convex functions that are “upside down”




- If  $f(x)$  is convex,  $-f(x)$  is concave.
- Some  $f(x)$  are **both** concave and convex
  - Example?

# Summary

- **Convex sets** are sets where a line segment between any two points is part of the set



- **Convex functions** are functions where the line segment between any two points is above the graph of the function




Break

# Unconstrained Optimization

---

# Unconstrained Minimization Problem


$$\underset{x}{\text{minimize}} f(x)$$



- Assumptions
  - $f$  is convex
  - No constraints on  $x$
- There are MANY methods for this kind of problem
  - Some are general, some exploit a specific structure of  $f$
- Usually decide what to use based on
  - Differentiability of  $f$
  - How you compute  $\nabla f(x)$
- We will cover several important methods common in robotics

# Review: Minimizing a simple function

- For a simple function, e.g.  $f(x) = x^2 - 4x$ , we can use calculus to find the minimum
- For an optimal point  $x^*$ ,  $\nabla f(x^*) = 0$
- So,
  1. Take the derivative of  $f(x)$
  2. Set it equal to 0
  3. Solve for  $x$
- Example for  $f(x) = x^2 - 4x$ 
  1.  $\nabla f(x) = 2x - 4$
  2.  $2x - 4 = 0$
  3.  $x = 2$  is the minimum




## What about a more complicated function?

- $f(x) = e^{0.5x+0.9} + e^{-0.5x^2-0.4} + 4x$ 
  1.  $\nabla f(x) = 0.5e^{0.5x+0.9} - xe^{-0.5x^2-0.4} + 4$
  2.  $0.5e^{0.5x+0.9} - xe^{-0.5x^2-0.4} + 4 = 0$
  3.  $x = ???$

**Problem:** No way to solve arbitrary equations using algebra!

minimize  $f(x)$

- assume  $f$  is convex
- assume  $x$  is unconstrained



\*Many more methods not covered!



# Descent Methods

---

# Unconstrained Minimization Methods



- Let  $p^*$  be the optimal value of  $f(x)$
- Let  $x^*$  be a value of  $x$  that produces  $p^*$ 
  - $p^* = f(x^*)$
- These methods produce a sequence of points:

$$x^{(k)} \in \text{dom } f, k = 0, 1, \dots$$

$$f(x^{(k)}) \rightarrow p^*$$

- Can interpret as iteratively finding an  $x^*$  that solves optimality condition:

$$\nabla f(x^*) = 0$$



# Descent methods

- We will cover two types of descent methods:
  - Gradient descent
  - Newton's method
    - Advantage: affine invariant
- Descent methods generate points with this property:

$$x^{(k+1)} = x^{(k)} + t^{(k)} \Delta x^{(k)} \quad \text{with } f(x^{(k+1)}) < f(x^{(k)})$$

Other notation:

$$x := x + t \Delta x$$

- $\Delta x$  is the **step**, or **search direction**
- $t$  is the **step size**, or **step length**

# General descent algorithm

Many ways  
to do these

given a starting point  $x \in \text{dom } f$ .

repeat


→ 1. Determine a descent direction  $\Delta x$ .

→ 2. *Line search*. Choose a step size  $t > 0$ .

→ 3. *Update*.  $x := x + t\Delta x$ .

until → stopping criterion is satisfied.

Input to the algorithm



# Gradient Descent

- Most common optimization algorithm
- Easy to implement, but may be slow to converge
- Descent direction:

$$\Delta x = -\nabla f(x)$$

- Termination condition:

$$\|\nabla f(x)\|_2 \leq \epsilon$$

e.g.  $\epsilon = 0.001$

# Gradient Descent: Step size

- Ideally, we would use *exact line search* to determine step size  $t$ :

$$t = \operatorname{argmin}_{t>0} f(x + t\Delta x)$$

- But this is slow to compute in general, so often use **backtracking line search**:

**given** a descent direction  $\Delta x$  for  $f$  at  $x \in \mathbf{dom} f$ ,  $\alpha \in (0, 0.5)$ ,  $\beta \in (0, 1)$ .  
 $t := 1$ .  
**while**  $f(x + t\Delta x) > f(x) + \alpha t \nabla f(x)^T \Delta x$ ,  $t := \beta t$ .


- I.e. decrease magnitude of step until you meet the stopping condition

# Gradient Descent Example

- Find the minimum of this function:

$$f(x_1, x_2) = 0.5(x_1^2 + 10x_2^2)$$

- Starting at  $x^{(0)} = (10, 1)$ , using exact line search




isocontours:  
levels of constant  
 $f(x)$  value


# Gradient Descent Example

- Find the minimum of this function:

$$f(x_1, x_2) = e^{x_1+3x_2-0.1} + e^{x_1-3x_2-0.1} + e^{-x_1-0.1}$$



backtracking line search



exact line search



# Problems with Gradient Descent

- Sensitive to the condition number of the Hessian
  - High condition number means very slow convergence
- Sensitive to the coordinates you use (not affine invariant)
  - Apply a linear transform to  $x$  and you may get different results!
- **Newton's method** overcomes these problems by using the Hessian of the function
  - For a price (the Hessian can be expensive to compute)

# Newton's method: Descent direction

- Determine a descent direction:



$$\Delta x_{\text{nt}} = -\nabla^2 f(x)^{-1} \nabla f(x)$$

- Why?
  - Let's approximate  $f(x)$  with a quadratic function (remember Taylor series):

$$\hat{f}(x + v) = f(x) + \nabla f(x)^T v + \frac{1}{2} v^T \nabla^2 f(x) v$$

- $x + \Delta x_{\text{nt}}$  solves the linearized optimality condition:

$$\nabla f(x + v) \approx \nabla \hat{f}(x + v) = \nabla f(x) + \nabla^2 f(x) v = 0$$



# Newton's method: Stopping criterion

- The **Newton decrement**  $\lambda(x)$  leads to the stopping criterion:

$$\lambda(x) = (\nabla f(x)^T \nabla^2 f(x)^{-1} \nabla f(x))^{1/2}$$

- $\lambda(x)$  is an estimate of the distance between  $f(x)$  and  $p^*$
- $\lambda(x)^2$  is the directional derivative in the direction of the Newton step:

$$\nabla f(x)^T \Delta x_{\text{nt}} = -\lambda(x)^2$$

- If the directional derivative is very close to 0,  $f(x)$  is not changing much in this direction
  - I.e. you're very close to the optimum
  - So, when  $\frac{\lambda(x)^2}{2}$  is below some small tolerance  $\epsilon$ , stop

# Newton's method

**given** a starting point  $x \in \text{dom } f$ , tolerance  $\epsilon > 0$ .

**repeat**

1. *Compute the Newton step and decrement.*

$$\Delta x_{\text{nt}} := -\nabla^2 f(x)^{-1} \nabla f(x); \quad \lambda^2 := \nabla f(x)^T \nabla^2 f(x)^{-1} \nabla f(x).$$

2. *Stopping criterion.* **quit** if  $\lambda^2/2 \leq \epsilon$ .



3. *Line search.* Choose step size  $t$  by backtracking line search.

4. *Update.*  $x := x + t\Delta x_{\text{nt}}$ .

# Newton's method example

- Find the optimum of this function:

$$f(x_1, x_2) = e^{x_1+3x_2-0.1} + e^{x_1-3x_2-0.1} + e^{-x_1-0.1}$$

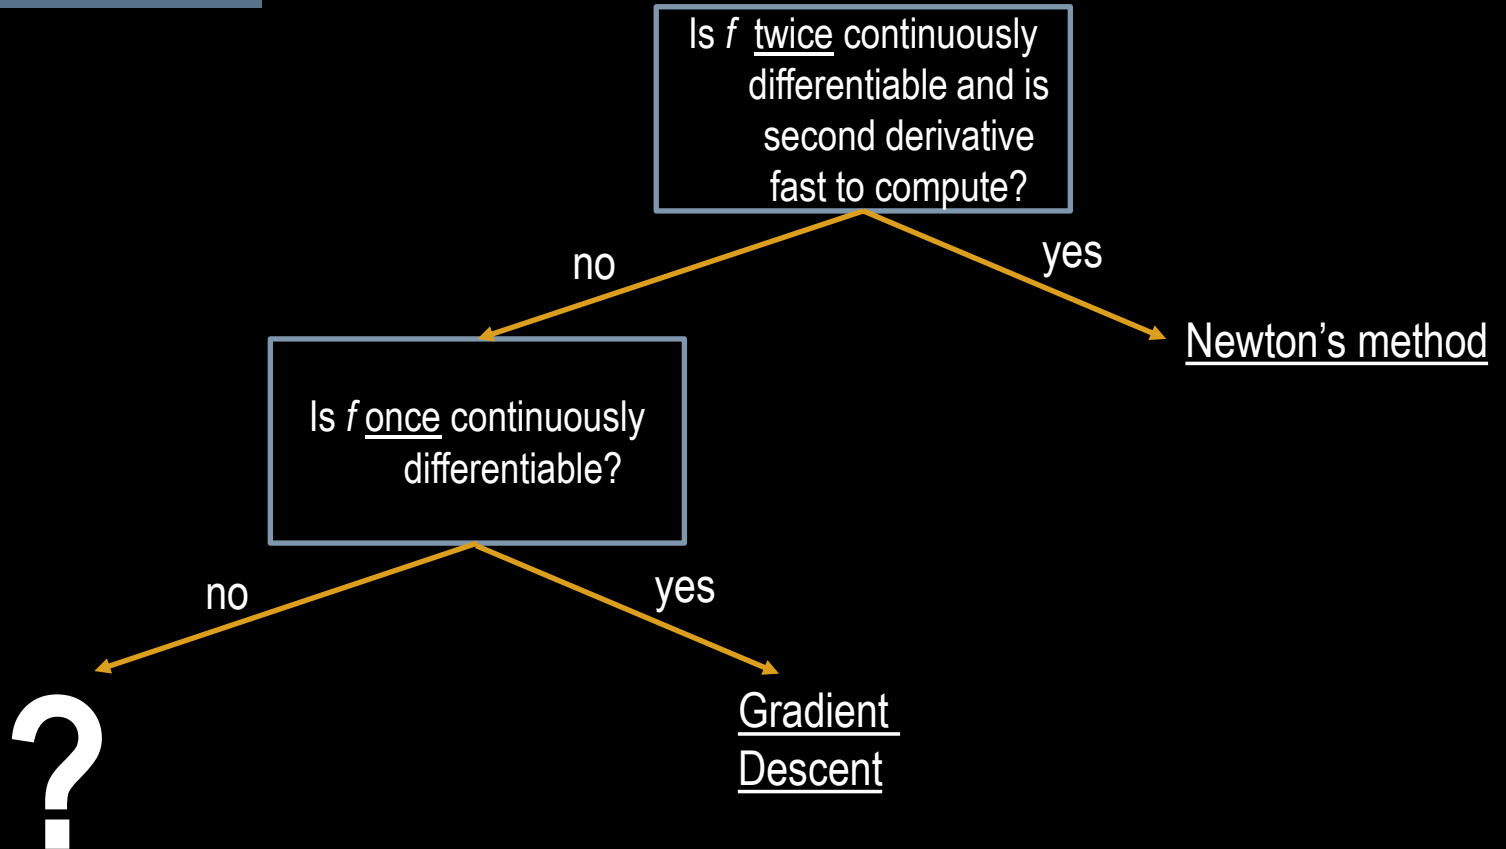


- Backtracking line search parameters:

$$\alpha = 0.1, \beta = 0.7$$

minimize  $f(x)$

- assume  $f$  is convex
- assume  $x$  is unconstrained



\*Many more methods not covered!


# Numerical Differentiation

---

# What about functions that you don't know analytically?

- So far  $f(x)$  is always represented analytically
- What if  $f(x)$  is this:

$x$  is actuator forces/torques



$f(x)$  outputs distance of tip to goal

“Optimization-based inverse model of Soft Robots with Contact Handling”  
Eulalie Coevoet, Adrien Escande, Christian Duriez



# Numerical Differentiation

- Need a way to differentiate when the function is not represented analytically
- Assume we can evaluate the function at any  $x$ 
  - E.g. by running some code like a simulation
- Recall standard derivative definition for  $f: \mathbf{R} \rightarrow \mathbf{R}$

$$Df(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- **Key idea:** Evaluate function at two points per dimension and estimate the derivative

# Numerical differentiation for univariate functions

1. Pick a small scalar  $h$
2. Use a **Finite Difference** method. Two common ones:

a) Newton's Difference Quotient

$$Df(x) \approx \frac{f(x+h) - f(x)}{h}$$

b) Symmetric Difference Quotient

$$Df(x) \approx \frac{f(x+h) - f(x-h)}{2h}$$

- There are other numerical methods which can give better estimates but use more function evaluations

# Numerical differentiation for multidimensional functions

- For  $f: \mathbf{R}^n \rightarrow \mathbf{R}^m$  we do the same thing to compute the Jacobian

- Recall:

$$Df(x)_{ij} = \frac{\partial f_i(x)}{\partial x_j}, \quad i = 1, \dots, m, \quad j = 1, \dots, n$$

index  $j$

- Let  $\delta(j, h) = [0, \dots, \overset{\downarrow}{h}, \dots 0]^T$

$$Df(x)_{ij} \approx \frac{f(x + \delta(j, h))_i - f(x)_i}{h}$$

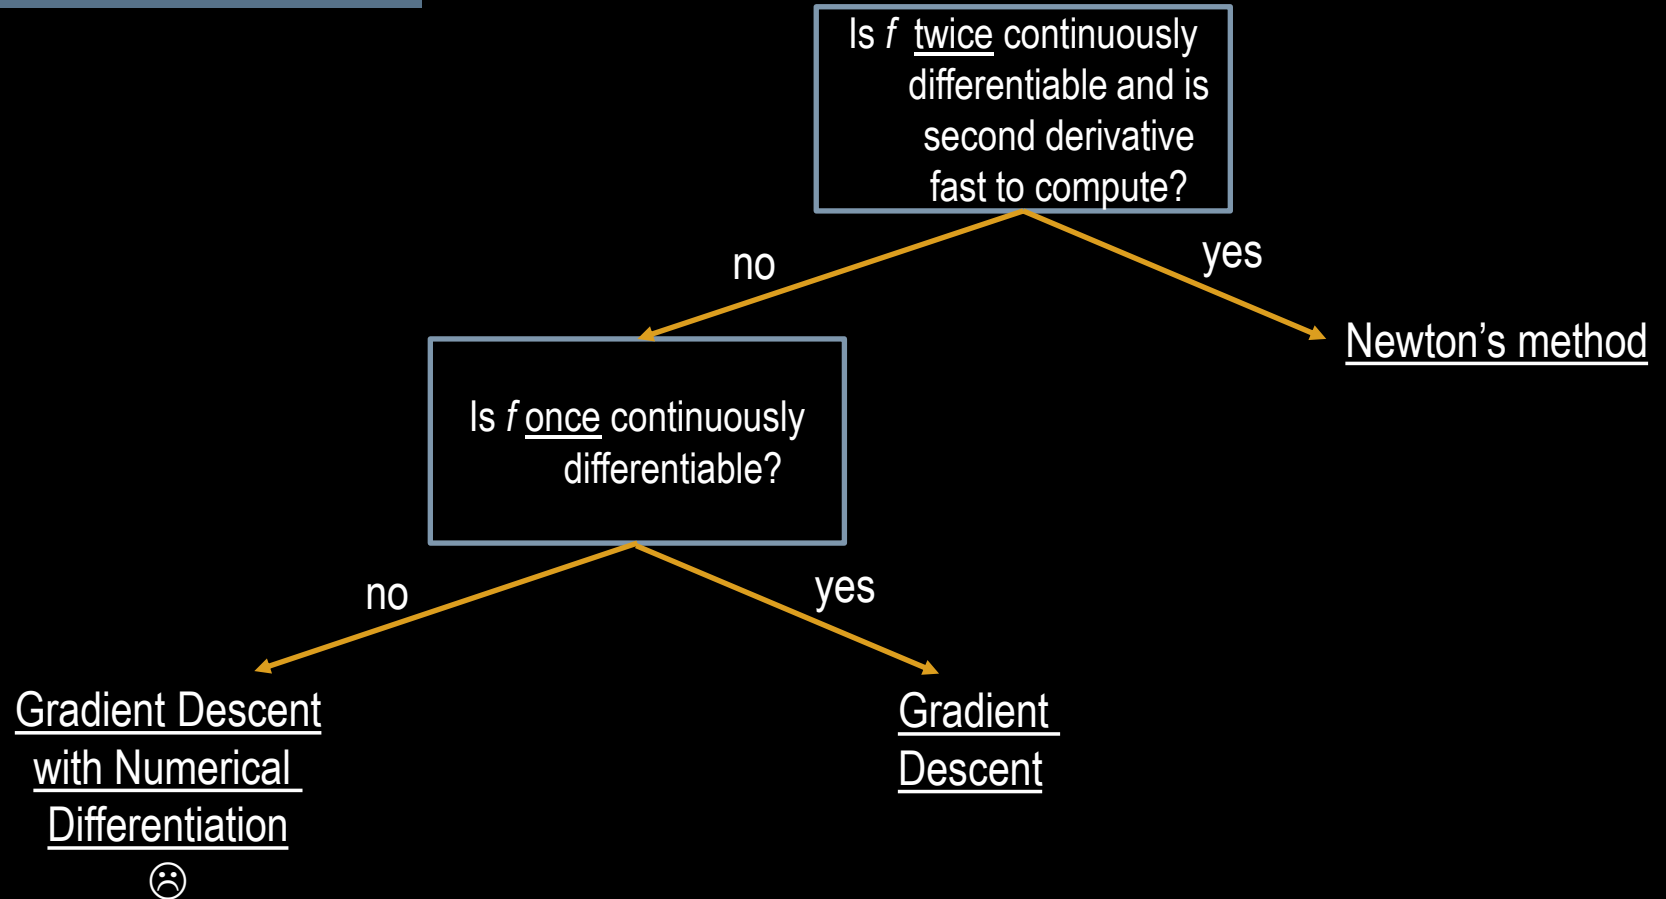
- Similar process for Symmetric Difference Quotient
- Thus we can use numerical differentiation to compute the gradient for gradient descent

# Limitations

- Choosing  $h$  well is difficult in general (it is function-dependent)
  - Many use a fixed  $h$  for simplicity
- Numerical methods can be very sensitive to the choice of  $h$
- There can be errors due to machine precision and floating point arithmetic

minimize  $f(x)$

- assume  $f$  is convex
- assume  $x$  is unconstrained



\*Many more methods not covered!

# Homework

- Reading from Optimization Book
  - Ch. 4.1-4.1.2, 4.3-4.3.1 (skip examples), 4.4-4.4.1 (only read first example in 4.4.1)
- Homework 4 due tonight
- Homework 5 posted tonight