



Searching for a Path

Last time...

- We saw how to frame optimization as a graph search problem



- But what if we need *a path* from a start to a goal that is feasible/optimal?
- Motion planning for robots requires solving this problem!

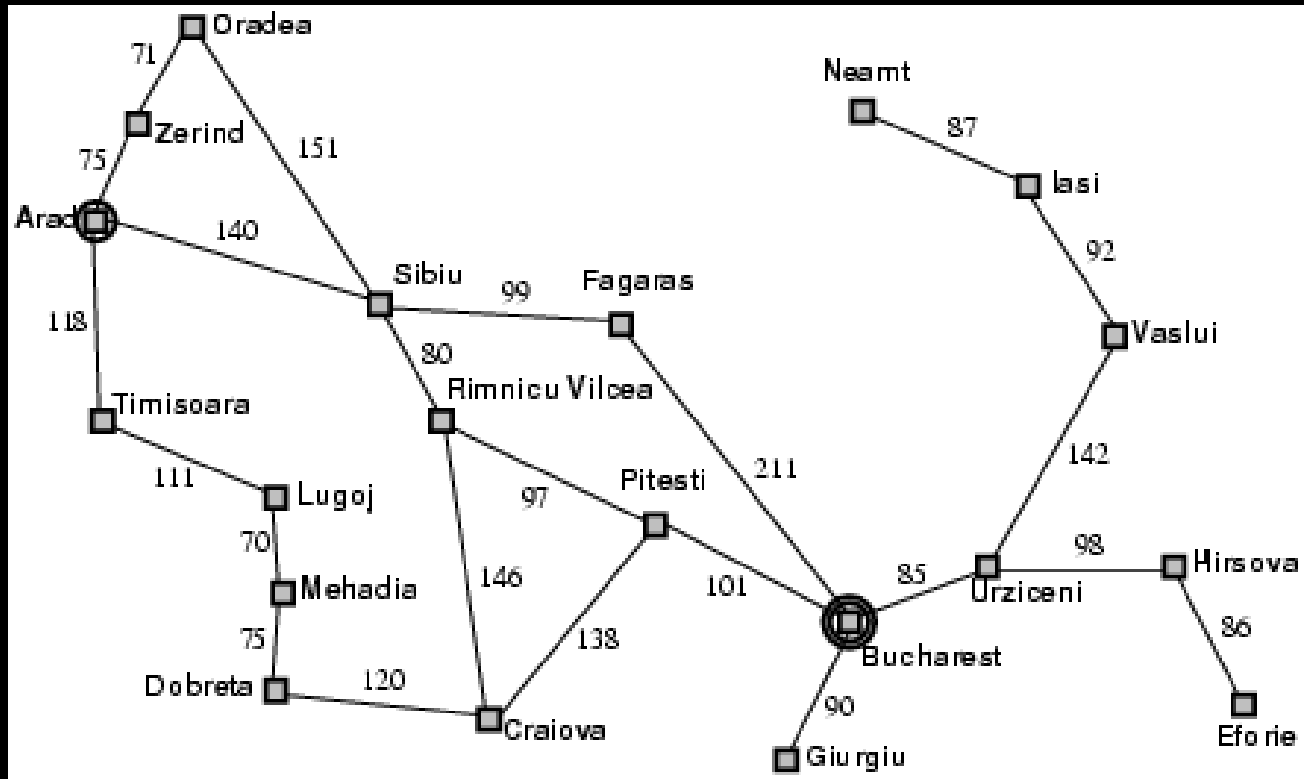
Outline

- Formulating a path search problem
 - Examples
 - General considerations for robotics
- Tree search algorithms
 - Uninformed search
 - Informed search

Formulating a path search problem

1. State Space
2. Successor Function
3. Actions
4. Action Cost
5. Goal Test

Example: Romanian roadmap



Example: A Romanian roadmap

1. State Space

- The set of cities

2. Successor Function

- A city's successors are cities directly connected to it (its neighbors)

3. Actions

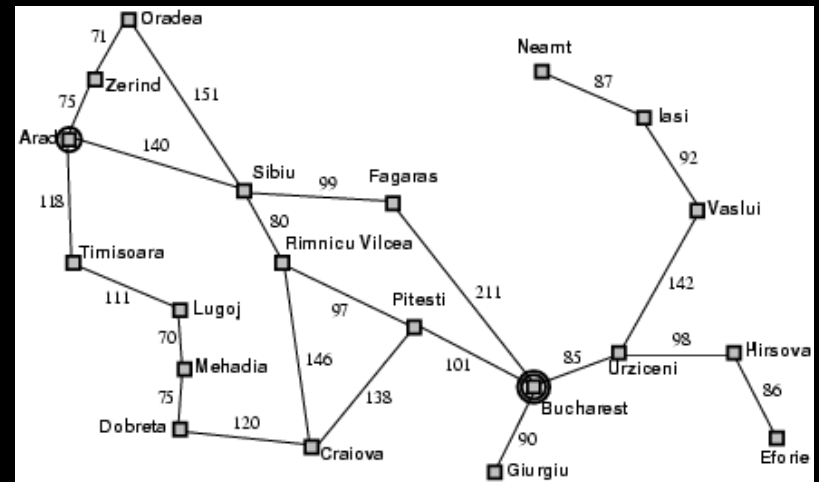
- Move to a neighboring city

4. Action Cost


- Distance between the cities

5. Goal Test

- Check if at goal city



Example : Point robot in a maze:



- Find a sequence of free cells that goes from start to goal

Point Robot Example

1. State Space

- The space of cells, usually in x,y coordinates

2. Successor Function

- A cell's successors are its free neighbors
- 4-connected vs. 8-connected

3. Actions



- Move to a neighboring cell

4. Action Cost


- Distance between cells traversed
- Are costs the same for 4 vs 8 connected?

5. Goal Test

- Check if at goal cell
- Multiple cells can be marked as goals




4-connected




8-connected

Formulating the problem: State space


- For motion planning in robotics, state space is often a grid
- There are many kinds of grids!




Cartesian Grid



Regular Grid



Rectilinear Grid






Curvilinear Grid

- The choice of grid (i.e. state space) is crucial to performance and accuracy
- Remember, the world is really continuous; these are all approximations

Formulating the problem: Actions

- Actions in motion planning are also often continuous
- For example, many ways to move between neighboring cells
- Usually pick a discrete action set *a priori*
- What are the tradeoffs in picking action sets?




from Knepper and
Mason, ICRA 2009



Formulating the problem: Successors

- These are largely determined by the action set
- Successors may not be known a priori
 - I.e. you have to try each action in your action set to see which cell you end in




Formulating the Problem: Action Cost

- Depends on what you're trying to optimize
 - Minimum Path Length: Cost is distance traversed when executing action
 - What is action cost for path smoothness?
- Sometimes we consider more than one criterion
 - Linear combination of cost functions (most common):
$$\text{Cost} = a_1C_1 + a_2C_2 + a_3C_3 \dots$$

Formulating the Problem: Goal Test

- Goals are most commonly specific cells you want to get to
- But they can be more abstract, too!
- Example Goals:
 - A state where X is visible
 - A state where the robot is contacting X
 - Topological goals



A topological goal example: go **right** around the obstacle (need whole path to evaluate if goal reached)

Finding a Path: Tree Search Algorithms

Tree Search Algorithms

```
function Tree-Search(problem, strategy)
    Root of search tree <- Initial state of the problem
    While 1
        If no nodes to expand
            return failure
        Choose a node  $n$  to expand according to strategy
        If  $n$  is a goal state
            return solution path //back-track from goal to
                                //start in the tree to get path
        Else
            NewNodes <- expand  $n$ 
            Add NewNodes as children of  $n$  in the search tree
```

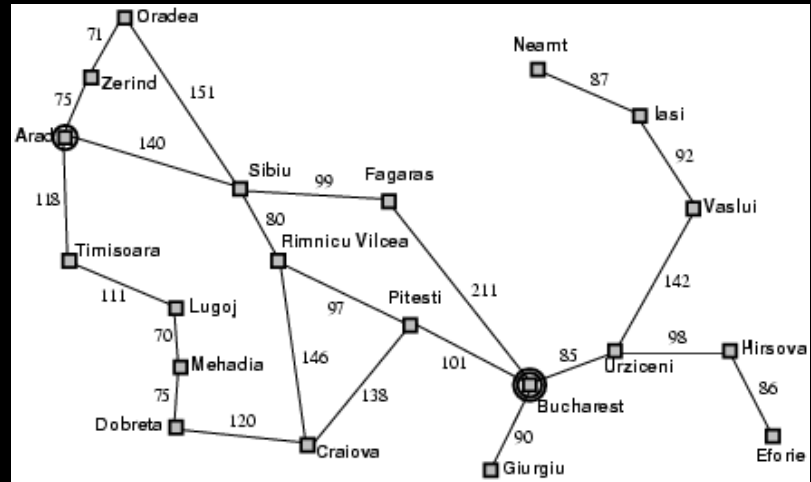
Strategies are evaluated in terms of

- **completeness**: does it always find a solution if one exists?
- **optimality**: does it always find a least-cost solution?
- Two types of complexity
 - **time complexity**: number of nodes visited
 - **space complexity**: maximum number of nodes in memory

Time and space complexity are measured in terms of


- b : maximum branching factor of the search tree (may $\rightarrow \infty$)
- d : depth of the least-cost solution
- m : maximum depth of the state space (may be ∞)

- What is b ? (worst case)
- What is d ?
- What is m ?



Tree Search Algorithm Implementation

- **Open list**: List of nodes we know about but haven't expanded yet
 - Implement as a queue, stack, or priority queue (depending on the algorithm)
- Need to avoid re-expanding the same state



- Solution: A **closed list** to track which nodes are already explored


Uninformed Search Strategies

Uninformed search strategies

- **Uninformed** search strategies use only the information available in the problem definition
- What does it mean to be uninformed?
 - You only know which states are connected by which actions. No additional information.
 - Later we'll talk about informed search, in which you can *estimate* which actions are likely to be better than others.

Breadth-first Search (BFS)

- **Main idea:** build search tree in layers
- Assumes all actions have equal cost
- Open list is a queue (e.g. `std::queue`), new nodes are inserted at the **back**




- Result: “Oldest” nodes are expanded first
- BFS finds the shortest path to the goal
- When would this strategy be very inefficient?

Properties of breadth-first search

- Complete?
 - Yes (if b is finite)
- Optimal?
 - Yes (if **cost = 1 per step**)
- Time?
 - $1+b+b^2+b^3+\dots+b^d = O(b^d)$
- Space?
 - $O(b^d)$ (keeps every node in memory)

Depth-first Search (DFS)

- **Main idea:** Go as deep as possible as fast as possible
- Assumes all actions have equal cost
- Open list is a stack, new nodes are inserted at the **front**
 - Use `std::stack`
 - Like a `std::queue`, except you can only add to the *front* of it, using `push()`
 - `std::stack` has nothing to do with the stack used in memory (i.e. stack vs. heap)




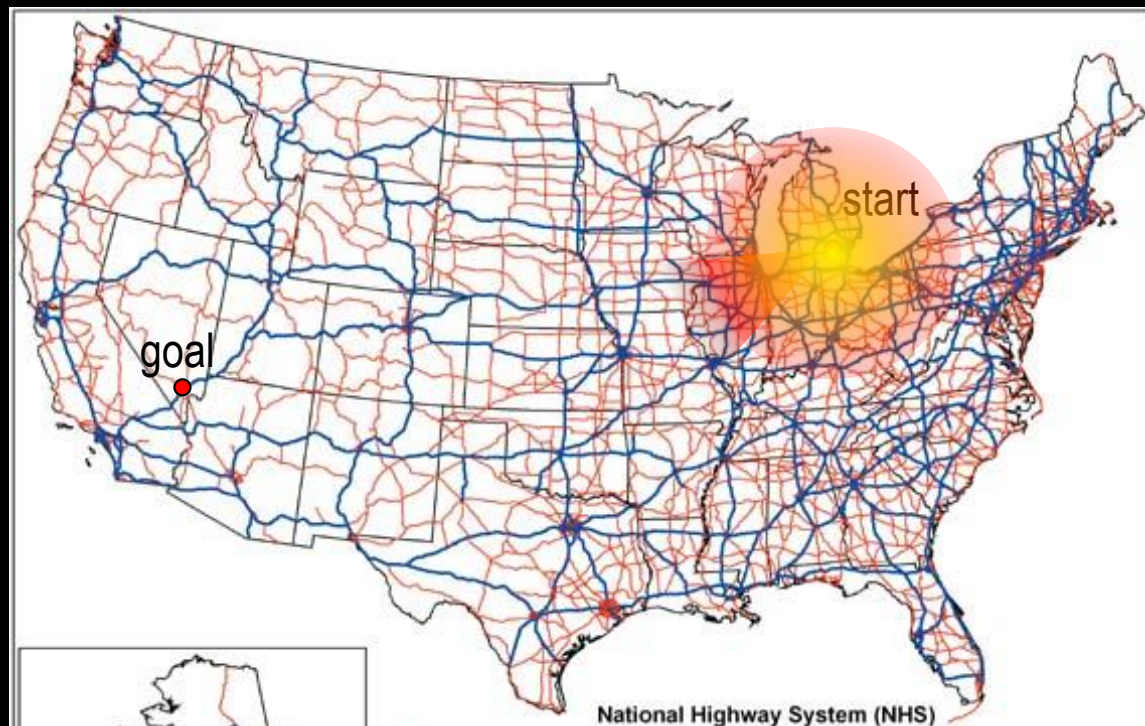
- Result: “Newest” nodes are expanded first
- DFS does NOT necessarily find the shortest path to the goal
- When would this strategy be very inefficient?

Properties of depth-first search


- Complete?
 - No: fails in infinite-depth spaces
 - Yes: in finite spaces
- Optimal?
 - No
- Time?
 - $O(b^m)$: (m is max depth of state space)
 - terrible if m is much larger than d
 - but if solutions are plentiful, may be much faster than breadth-first
- Space?
 - $O(bm)$

Informed Search Strategies

Search types



Uninformed search



Informed search

We'll need a new tool: `std::priority_queue`

- Similar to `std::queue`, but instead of the order being in order of insertion (last-in-first-out), in `std::priority_queue`, elements are sorted by their *priority*
- For simple types, it uses “<” to sort:

```
std::priority_queue<int> pq_int;
```

```
pq_int.push(1);  
pq_int.push(7);  
pq_int.push(5);  
pq_int.push(3);  
pq_int.push(285);
```

```
while (!pq_int.empty()) {  
    int p = pq_int.top();  
    pq_int.pop();  
    cout << p << " ";  
}  
//prints 285 7 5 3 1
```

This is how you
get the first
element

Same as for
`std::queue`

- Note that it re-sorts every time you push an element!

We'll need a new tool: `std::priority_queue`

- For more complex types, need to define "<", e.g:

```
class Node {
public:
    int priority; //bad style: these variables should be private
    string name;
    Node(string name, int priority) : name(name), priority(priority){}
};

bool operator<(const Node& p1, const Node& p2)
{
    return p1.priority < p2.priority;
}

int main()
{
    priority_queue<Node> pq;
    pq.push(Node("N1",21));
    pq.push(Node("N2",19));
    pq.push(Node("N3",47));
    pq.push(Node("N4",23));


    while (!pq.empty()) {
        Node n = pq.top();
        pq.pop();
        cout << n.name << " " << n.priority << "\n";
    }
    return 0;
}
```

Output:

```
N3 47
N4 23
N1 21
N2 19
```


Best-first Search

- **Main idea:** Use Heuristic function $h(n)$ to estimate each node's distance to goal, expand node with minimum $h(n)$
- Open list is a *priority queue*, nodes are sorted according to $h(n)$ (ascending)
- Result: Works great if heuristic is a good estimate
- Does not necessarily find least-cost path
- When would this strategy be inefficient?



A * Evaluation functions

- **Main idea:** Select nodes based on estimated path to goal
- Evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ = cost so far to reach n (**cost-to-come**)
 - $h(n)$ = estimated cost from n to goal (**cost-to-go**)
 - $f(n)$ = **estimated total cost** of path through n to goal



A* Search

- Open list is a *priority queue*, nodes are sorted according to $f(n)$ (ascending)
- $g(n)$ is sum of edge costs from root node to n



A* Search

- IMPORTANT RESULT: If $h(n)$ is *admissible*, A* will find the least-cost path!
- **Admissibility**: $h(n)$ must *never overestimate* the true cost to reach the goal from n
 - $h(n) \leq h^*(n)$, where $h^*(n)$ is the true cost to reach goal from n
 - $h(n) \geq 0$ (so $h(G) = 0$ for goals G)
- “Inflating” the heuristic may give you faster search, but least-cost path is not guaranteed

Properties of A^* (w/ admissible heuristic)

- Complete?
 - Yes (unless there are infinitely many nodes with $f \leq f(G)$)
- Optimal?
 - Yes
- Time?
 - Exponential, approximately $O(b^d)$ in the worst case
- Space?
 - $O(b^m)$ Keeps all nodes in memory

Heuristic Consistency

- A heuristic is **consistent** (also known as monotonic) when it obeys

$$h(n_1) \leq C(n_1, n_2) + h(n_2)$$

$$h(G) = 0 \text{ for goals } G$$

for adjacent n_1, n_2 and with edge cost $C(n_1, n_2)$

- Ensures that it is impossible to decrease f by extending a path to include a neighboring node
- A consistent heuristic is always admissible
- If a heuristic is consistent, we can use a closed list
- If not consistent, we cannot use a closed list and guarantee optimality

Summary

- A search problem is defined by
 1. State Space
 2. Successor Function
 3. Actions
 4. Action Cost
 5. Goal Test
- Search types:
 - Uninformed: Uses only problem definition, no additional information
 - Informed: Relies on a heuristic to estimate cost to reach the goal
- A* is a strong approach
 - My default method for search problems

Homework

- Homework 6